

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier,
RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score
import matplotlib.pyplot as plt
import seaborn as sns

```

Load the dataset

```
df = pd.read_csv('heart_statlog_cleveland_hungary_final.csv')
```

Display the first few rows and basic information about the dataset

```

print(df.head())
print("\nDataset Info:")
print(df.info())

```

Split the data into features (X) and target (y)

```

X = df.drop('target', axis=1)
y = df['target']

```

Split the data into training (80%) and test (20%) sets

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Scale the features

```

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```
print("\nShape of training set:", X_train_scaled.shape)
```

```
print("Shape of test set:", X_test_scaled.shape)
```

```
print("Dataset preparation completed.")
```

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar \
0	40	1	2	140	289	
1	49	0	3	160	180	
2	37	1	2	130	283	
3	48	0	4	138	214	
4	54	1	3	150	195	

```

0
    resting ecg  max heart rate  exercise angina  oldpeak  ST slope
target
0              0             172                0      0.0        1
0
1              0             156                0      1.0        2
1
2              1              98                0      0.0        1
0
3              0             108                1      1.5        2
1
4              0             122                0      0.0        1
0

```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1190 entries, 0 to 1189
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1190 non-null	int64
1	sex	1190 non-null	int64
2	chest pain type	1190 non-null	int64
3	resting bp s	1190 non-null	int64
4	cholesterol	1190 non-null	int64
5	fasting blood sugar	1190 non-null	int64
6	resting ecg	1190 non-null	int64
7	max heart rate	1190 non-null	int64
8	exercise angina	1190 non-null	int64
9	oldpeak	1190 non-null	float64
10	ST slope	1190 non-null	int64
11	target	1190 non-null	int64

```
dtypes: float64(1), int64(11)
```

```
memory usage: 111.7 KB
```

```
None
```

```
Shape of training set: (952, 11)
```

```
Shape of test set: (238, 11)
```

```
Dataset preparation completed.
```

```

from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier,
RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score
from scipy.stats import randint, uniform

```

```
# Define the models
```

```

svm = SVC(random_state=42, probability=True) # Set probability=True
gbm = GradientBoostingClassifier(random_state=42)
rf = RandomForestClassifier(random_state=42)

# Define parameter distributions for RandomizedSearchCV
svm_params = {
    'C': uniform(0.1, 10),
    'kernel': ['rbf', 'linear', 'poly', 'sigmoid'],
    'gamma': uniform(0.01, 1)
}

gbm_params = {
    'n_estimators': randint(50, 300),
    'learning_rate': uniform(0.01, 0.3),
    'max_depth': randint(3, 10),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10)
}

rf_params = {
    'n_estimators': randint(50, 300),
    'max_depth': randint(3, 15),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10),
    'max_features': ['auto', 'sqrt', 'log2']
}

# Function to perform RandomizedSearchCV and return best model
def tune_model(model, params, X_train, y_train):
    rs_cv = RandomizedSearchCV(model, params, n_iter=50, cv=5,
n_jobs=-1, random_state=42, scoring='roc_auc')
    rs_cv.fit(X_train, y_train)
    return rs_cv.best_estimator_

print("Models and parameters defined successfully.")

Models and parameters defined successfully.

# Tune models
print("Tuning SVM...")
best_svm = tune_model(svm, svm_params, X_train_scaled, y_train)

print("Tuning GBM...")
best_gbm = tune_model(gbm, gbm_params, X_train_scaled, y_train)

print("Tuning Random Forest...")
best_rf = tune_model(rf, rf_params, X_train_scaled, y_train)

print("Model tuning completed.")

```

```
Tuning SVM...
Tuning GBM...
Tuning Random Forest...
Model tuning completed.
```

```
/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/
model_selection/_validation.py:540: FitFailedWarning:
80 fits failed out of a total of 250.
The score on these train-test partitions for these parameters will be
set to nan.
If these failures are not expected, you can try to debug them by
setting error_score='raise'.
```

```
Below are more details about the failures:
```

```
-----
-----
80 fits failed with the following error:
Traceback (most recent call last):
  File
"/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/model_s
election/_validation.py", line 888, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File
"/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/base.py
", line 1466, in wrapper
    estimator._validate_params()
  File
"/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/base.py
", line 666, in _validate_params
    validate_parameter_constraints(
  File
"/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/utils/_
param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestClassifier must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
/workspaces/Team-7/.venv/lib/python3.12/site-packages/sklearn/model_se
lection/_search.py:1103: UserWarning: One or more of the test scores
are non-finite: [      nan  0.92259001  0.93258719  0.93225552
nan      nan
0.93839674      nan      nan      nan  0.93140273      nan
0.9128019  0.93078244  0.93515735  0.92980213  0.91887379      nan
0.91401772  0.92972532  0.92579938  0.93692581  0.9278099  0.94347259
0.93647097  0.92958122  0.9418757      nan      nan  0.94149868
      nan  0.92686167      nan  0.92405137  0.93646385  0.93205586
0.91295694  0.92106995      nan  0.93018049      nan  0.93187963
```

```
0.94583186 0.93108955 nan 0.92635855 0.93106527 0.94112334
nan 0.92958408]
warnings.warn(
```

```
# Function to evaluate model
```

```
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision, recall, f1, _ = precision_recall_fscore_support(y_test,
y_pred, average='weighted')
    auc_roc = roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])
    return accuracy, precision, recall, f1, auc_roc
```

```
# Evaluate models
```

```
models = {'SVM': best_svm, 'GBM': best_gbm, 'Random Forest': best_rf}
results = {}
```

```
for name, model in models.items():
    print(f"\nEvaluating {name}...")
    accuracy, precision, recall, f1, auc_roc = evaluate_model(model,
X_test_scaled, y_test)
    results[name] = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-score': f1,
        'AUC-ROC': auc_roc
    }
    print(f"{name} Results:")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-score: {f1:.4f}")
    print(f"AUC-ROC: {auc_roc:.4f}")
```

```
# Create a summary table
```

```
summary_df = pd.DataFrame(results).T
print("\nSummary Table:")
print(summary_df)
```

```
# Plot the results
```

```
plt.figure(figsize=(12, 6))
summary_df.plot(kind='bar', ax=plt.gca())
plt.title('Model Comparison')
plt.xlabel('Models')
plt.ylabel('Scores')
plt.legend(title='Metrics', bbox_to_anchor=(1.05, 1), loc='upper
left')
plt.tight_layout()
plt.savefig('model_comparison.png')
```

```
print("Model comparison plot saved as 'model_comparison.png'")
```

```
# Print best parameters for each model
```

```
print("\nBest Parameters:")
```

```
print("SVM:", best_svm.get_params())
```

```
print("GBM:", best_gbm.get_params())
```

```
print("Random Forest:", best_rf.get_params())
```

Evaluating SVM...

SVM Results:

Accuracy: 0.8992

Precision: 0.9023

Recall: 0.8992

F1-score: 0.8984

AUC-ROC: 0.9700

Evaluating GBM...

GBM Results:

Accuracy: 0.9496

Precision: 0.9496

Recall: 0.9496

F1-score: 0.9495

AUC-ROC: 0.9799

Evaluating Random Forest...

Random Forest Results:

Accuracy: 0.9370

Precision: 0.9375

Recall: 0.9370

F1-score: 0.9368

AUC-ROC: 0.9681

Summary Table:

	Accuracy	Precision	Recall	F1-score	AUC-ROC
SVM	0.899160	0.902322	0.899160	0.898373	0.970036
GBM	0.949580	0.949622	0.949580	0.949533	0.979882
Random Forest	0.936975	0.937490	0.936975	0.936811	0.968110

Model comparison plot saved as 'model_comparison.png'

Best Parameters:

SVM: {'C': np.float64(7.390071680409873), 'break_ties': False, 'cache_size': 200, 'class_weight': None, 'coef0': 0.0, 'decision_function_shape': 'ovr', 'degree': 3, 'gamma': np.float64(0.7812703466859457), 'kernel': 'rbf', 'max_iter': -1, 'probability': True, 'random_state': 42, 'shrinking': True, 'tol': 0.001, 'verbose': False}

GBM: {'ccp_alpha': 0.0, 'criterion': 'friedman_mse', 'init': None, 'learning_rate': np.float64(0.2921569793468812), 'loss': 'log_loss', 'max_depth': 9, 'max_features': None, 'max_leaf_nodes': None,

```
'min_impurity_decrease': 0.0, 'min_samples_leaf': 2,  
'min_samples_split': 9, 'min_weight_fraction_leaf': 0.0,  
'n_estimators': 178, 'n_iter_no_change': None, 'random_state': 42,  
'subsample': 1.0, 'tol': 0.0001, 'validation_fraction': 0.1,  
'verbose': 0, 'warm_start': False}  
Random Forest: {'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight':  
None, 'criterion': 'gini', 'max_depth': 13, 'max_features': 'log2',  
'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease':  
0.0, 'min_samples_leaf': 1, 'min_samples_split': 7,  
'min_weight_fraction_leaf': 0.0, 'monotonic_cst': None,  
'n_estimators': 73, 'n_jobs': None, 'oob_score': False,  
'random_state': 42, 'verbose': 0, 'warm_start': False}
```

