

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier,
RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
df = pd.read_csv(r'C:\Users\saket\Downloads\archive (4)\
heart_statlog_cleveland_hungary_final.csv')

# Display the first few rows and basic information about the dataset
print(df.head())
print("\nDataset Info:")
print(df.info())

# Split the data into features (X) and target (y)
X = df.drop('target', axis=1)
y = df['target']

# Split the data into training (80%) and test (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("\nShape of training set:", X_train_scaled.shape)
print("Shape of test set:", X_test_scaled.shape)

print("Dataset preparation completed.")

```

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar \
0	40	1	2	140	289	
1	49	0	3	160	180	
2	37	1	2	130	283	
3	48	0	4	138	214	

```

0
4  54    1          3          150          195
0
    resting ecg  max heart rate  exercise angina  oldpeak  ST slope
target
0              0             172              0      0.0      1
0
1              0             156              0      1.0      2
1
2              1              98              0      0.0      1
0
3              0             108              1      1.5      2
1
4              0             122              0      0.0      1
0

```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1190 entries, 0 to 1189

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	age	1190 non-null	int64
1	sex	1190 non-null	int64
2	chest pain type	1190 non-null	int64
3	resting bp s	1190 non-null	int64
4	cholesterol	1190 non-null	int64
5	fasting blood sugar	1190 non-null	int64
6	resting ecg	1190 non-null	int64
7	max heart rate	1190 non-null	int64
8	exercise angina	1190 non-null	int64
9	oldpeak	1190 non-null	float64
10	ST slope	1190 non-null	int64
11	target	1190 non-null	int64

dtypes: float64(1), int64(11)

memory usage: 111.7 KB

None

Shape of training set: (952, 11)

Shape of test set: (238, 11)

Dataset preparation completed.

Define the models

```
svm = SVC(random_state=42, probability=True) # Set probability=True
```

```
gbm = GradientBoostingClassifier(random_state=42)
```

```
rf = RandomForestClassifier(random_state=42)
```

Define parameter distributions for RandomizedSearchCV

```
svm_params = {
```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier,
RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset
df = pd.read_csv(r'C:\Users\saket\Downloads\archive (4)\
heart_statlog_cleveland_hungary_final.csv')

# Display the first few rows and basic information about the dataset
print(df.head())
print("\nDataset Info:")
print(df.info())

# Split the data into features (X) and target (y)
X = df.drop('target', axis=1)
y = df['target']

# Split the data into training (80%) and test (20%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("\nShape of training set:", X_train_scaled.shape)
print("Shape of test set:", X_test_scaled.shape)

print("Dataset preparation completed.")

```

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar \
0	40	1	2	140	289	
1	49	0	3	160	180	
2	37	1	2	130	283	
3	48	0	4	138	214	

```

b
4  54    1          3          150          195
0

```

	resting ecg	max heart rate	exercise angina	oldpeak	ST slope	target
0	0	172	0	0.0	1	
0						
1	0	156	0	1.0	2	
1						
2	1	98	0	0.0	1	
0						
3	0	108	1	1.5	2	
1						
4	0	122	0	0.0	1	
0						

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1190 entries, 0 to 1189
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	age	1190 non-null	int64
1	sex	1190 non-null	int64
2	chest pain type	1190 non-null	int64
3	resting bp s	1190 non-null	int64
4	cholesterol	1190 non-null	int64
5	fasting blood sugar	1190 non-null	int64
6	resting ecg	1190 non-null	int64
7	max heart rate	1190 non-null	int64
8	exercise angina	1190 non-null	int64
9	oldpeak	1190 non-null	float64
10	ST slope	1190 non-null	int64
11	target	1190 non-null	int64

```
dtypes: float64(1), int64(11)
```

```
memory usage: 111.7 KB
```

```
None
```

```
Shape of training set: (952, 11)
```

```
Shape of test set: (238, 11)
```

```
Dataset preparation completed.
```

```
# Define the models
```

```
svm = SVC(random_state=42, probability=True) # Set probability=True
```

```
gbm = GradientBoostingClassifier(random_state=42)
```

```
rf = RandomForestClassifier(random_state=42)
```

```
# Define parameter distributions for RandomizedSearchCV
```

```
svm_params = {
```

```

    'C': uniform(0.1, 10),
    'kernel': ['rbf', 'linear', 'poly', 'sigmoid'],
    'gamma': uniform(0.01, 1)
}

gbm_params = {
    'n_estimators': randint(50, 300),
    'learning_rate': uniform(0.01, 0.3),
    'max_depth': randint(3, 10),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10)
}

rf_params = {
    'n_estimators': randint(50, 300),
    'max_depth': randint(3, 15),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10),
    'max_features': ['auto', 'sqrt', 'log2']
}

# Function to perform RandomizedSearchCV and return best model
def tune_model(model, params, X_train, y_train):
    rs_cv = RandomizedSearchCV(model, params, n_iter=50, cv=5,
n_jobs=-1, random_state=42, scoring='roc_auc')
    rs_cv.fit(X_train, y_train)
    return rs_cv.best_estimator_

# Tune models
print("Tuning SVM...")
best_svm = tune_model(svm, svm_params, X_train_scaled, y_train)

print("Tuning GBM...")
best_gbm = tune_model(gbm, gbm_params, X_train_scaled, y_train)

print("Tuning Random Forest...")
best_rf = tune_model(rf, rf_params, X_train_scaled, y_train)

# Function to evaluate model
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision, recall, f1, _ = precision_recall_fscore_support(y_test,
y_pred, average='weighted')
    auc_roc = roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])
    return accuracy, precision, recall, f1, auc_roc

# Evaluate models
models = {'SVM': best_svm, 'GBM': best_gbm, 'Random Forest': best_rf}
results = {}

```

```

for name, model in models.items():
    print(f"\nEvaluating {name}...")
    accuracy, precision, recall, f1, auc_roc = evaluate_model(model,
X test scaled, y test)
    results[name] = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-score': f1,
        'AUC-ROC': auc_roc
    }
    print(f"{name} Results:")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-score: {f1:.4f}")
    print(f"AUC-ROC: {auc_roc:.4f}")

# Create a summary table
summary_df = pd.DataFrame(results).T
print("\nSummary Table:")
print(summary_df)

# Plot the results
plt.figure(figsize=(12, 6))
summary_df.plot(kind='bar', ax=plt.gca())
plt.title('Model Comparison')
plt.xlabel('Models')
plt.ylabel('Scores')
plt.legend(title='Metrics', bbox_to_anchor=(1.05, 1), loc='upper
left')
plt.tight_layout()
plt.savefig('model_comparison.png')
print("Model comparison plot saved as 'model_comparison.png'")

# Print best parameters for each model
print("\nBest Parameters:")
print("SVM:", best_svm.get_params())
print("GBM:", best_gbm.get_params())
print("Random Forest:", best_rf.get_params())

Tuning SVM...

C:\Users\saket\anaconda3\Lib\site-packages\joblib\externals\loky\
process_executor.py:700: UserWarning: A worker stopped while some jobs
were given to the executor. This can be caused by a too short worker
timeout or by a memory leak.
    warnings.warn(

```

Tuning GBM...

Tuning Random Forest...

```
C:\Users\saket\anaconda3\Lib\site-packages\sklearn\ensemble\
forest.py:424: FutureWarning: `max_features='auto'` has been
deprecated in 1.1 and will be removed in 1.3. To keep the past
behaviour, explicitly set `max_features='sqrt'` or remove this
parameter as it is also the default value for RandomForestClassifiers
and ExtraTreesClassifiers.
  warn(
```

Evaluating SVM...

SVM Results:

Accuracy: 0.8992
Precision: 0.9023
Recall: 0.8992
F1-score: 0.8984
AUC-ROC: 0.9700

Evaluating GBM...

GBM Results:

Accuracy: 0.9412
Precision: 0.9412
Recall: 0.9412
F1-score: 0.9411
AUC-ROC: 0.9786

Evaluating Random Forest...

Random Forest Results:

Accuracy: 0.9412
Precision: 0.9412
Recall: 0.9412
F1-score: 0.9411
AUC-ROC: 0.9691

Summary Table:

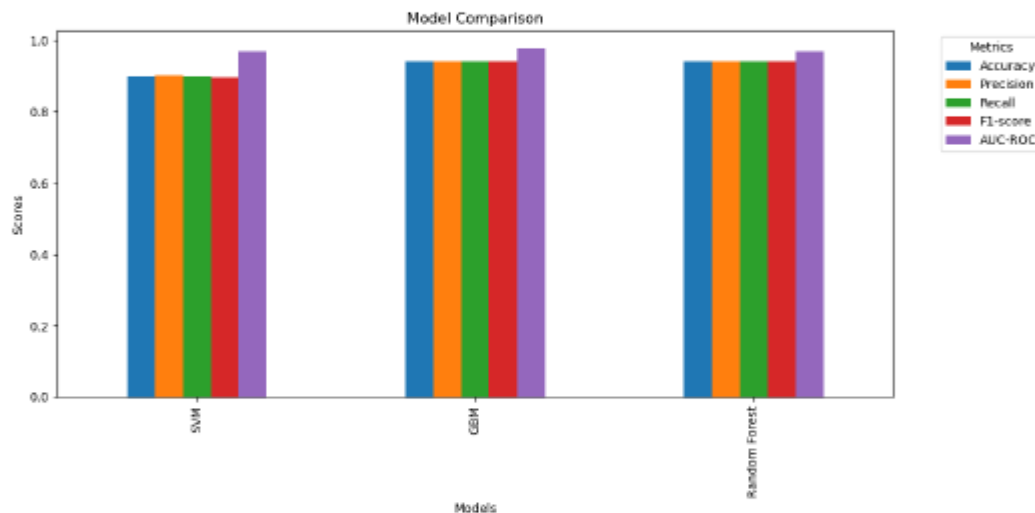
	Accuracy	Precision	Recall	F1-score	AUC-ROC
SVM	0.899160	0.902322	0.899160	0.898373	0.970036
GBM	0.941176	0.941202	0.941176	0.941122	0.978597
Random Forest	0.941176	0.941202	0.941176	0.941122	0.969109

Model comparison plot saved as 'model_comparison.png'

Best Parameters:

SVM: {'C': 7.390071680409873, 'break_ties': False, 'cache_size': 200, 'class_weight': None, 'coef0': 0.0, 'decision_function_shape': 'ovr', 'degree': 3, 'gamma': 0.7812703466859457, 'kernel': 'rbf', 'max_iter': -1, 'probability': True, 'random_state': 42, 'shrinking': True, 'tol': 0.001, 'verbose': False}
GBM: {'ccp_alpha': 0.0, 'criterion': 'friedman_mse', 'init': None,


```
|'learning rate': 0.2921569793468812, 'loss': 'log_loss', 'max_depth':
9, 'max_features': None, 'max_leaf_nodes': None,
'min_impurity_decrease': 0.0, 'min_samples_leaf': 2,
'min_samples_split': 9, 'min_weight_fraction_leaf': 0.0,
'n_estimators': 178, 'n_iter_no_change': None, 'random_state': 42,
'subsample': 1.0, 'tol': 0.0001, 'validation_fraction': 0.1,
'verbose': 0, 'warm_start': False}
Random Forest: {'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight':
None, 'criterion': 'gini', 'max_depth': 10, 'max_features': 'auto',
'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease':
0.0, 'min_samples_leaf': 1, 'min_samples_split': 3,
'min_weight_fraction_leaf': 0.0, 'n_estimators': 179, 'n_jobs': None,
'oob_score': False, 'random_state': 42, 'verbose': 0, 'warm_start':
False}
```



```
from sklearn.svm import SVC
from sklearn.ensemble import GradientBoostingClassifier,
RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score
from scipy.stats import randint, uniform

# Define the models
svm = SVC(random_state=42, probability=True) # Set probability=True
gbm = GradientBoostingClassifier(random_state=42)
rf = RandomForestClassifier(random_state=42)

# Define parameter distributions for RandomizedSearchCV
svm_params = {
```



```

    'C': uniform(0.1, 10),
    'kernel': ['rbf', 'linear', 'poly', 'sigmoid'],
    'gamma': uniform(0.01, 1)
}

gbm_params = {
    'n_estimators': randint(50, 300),
    'learning_rate': uniform(0.01, 0.3),
    'max_depth': randint(3, 10),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10)
}

rf_params = {
    'n_estimators': randint(50, 300),
    'max_depth': randint(3, 15),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 10),
    'max_features': ['auto', 'sqrt', 'log2']
}

# Function to perform RandomizedSearchCV and return best model
def tune_model(model, params, X_train, y_train):
    rs_cv = RandomizedSearchCV(model, params, n_iter=50, cv=5,
n_jobs=-1, random_state=42, scoring='roc_auc')
    rs_cv.fit(X_train, y_train)
    return rs_cv.best_estimator_

print("Models and parameters defined successfully.")

Models and parameters defined successfully.

# Tune models
print("Tuning SVM...")
best_svm = tune_model(svm, svm_params, X_train_scaled, y_train)

print("Tuning GBM...")
best_gbm = tune_model(gbm, gbm_params, X_train_scaled, y_train)

print("Tuning Random Forest...")
best_rf = tune_model(rf, rf_params, X_train_scaled, y_train)

print("Model tuning completed.")

Tuning SVM...
Tuning GBM...
Tuning Random Forest...

C:\Users\saket\anaconda3\Lib\site-packages\sklearn\ensemble\
forest.py:424: FutureWarning: `max_features='auto'` has been
deprecated in 1.1 and will be removed in 1.3. To keep the past

```

behaviour, explicitly set `max features='sqrt'` or remove this parameter as it is also the default value for RandomForestClassifiers and ExtraTreesClassifiers.

```
warn()
```

Model tuning completed.

```
# Function to evaluate model
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision, recall, f1, _ = precision_recall_fscore_support(y_test,
y_pred, average='weighted')
    auc_roc = roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])
    return accuracy, precision, recall, f1, auc_roc

# Evaluate models
models = {'SVM': best_svm, 'GBM': best_gbm, 'Random Forest': best_rf}
results = {}

for name, model in models.items():
    print(f"\nEvaluating {name}...")
    accuracy, precision, recall, f1, auc_roc = evaluate_model(model,
X_test_scaled, y_test)
    results[name] = {
        'Accuracy': accuracy,
        'Precision': precision,
        'Recall': recall,
        'F1-score': f1,
        'AUC-ROC': auc_roc
    }
    print(f"{name} Results:")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1-score: {f1:.4f}")
    print(f"AUC-ROC: {auc_roc:.4f}")

# Create a summary table
summary_df = pd.DataFrame(results).T
print("\nSummary Table:")
print(summary_df)

# Plot the results
plt.figure(figsize=(12, 6))
summary_df.plot(kind='bar', ax=plt.gca())
plt.title('Model Comparison')
plt.xlabel('Models')
plt.ylabel('Scores')
plt.legend(title='Metrics', bbox_to_anchor=(1.05, 1), loc='upper
```

```

left')
plt.tight_layout()
plt.savefig('model_comparison.png')
print("Model comparison plot saved as 'model_comparison.png'")

# Print best parameters for each model
print("\nBest Parameters:")
print("SVM:", best_svm.get_params())
print("GBM:", best_gbm.get_params())
print("Random Forest:", best_rf.get_params())

```

Evaluating SVM...

SVM Results:

Accuracy: 0.8992
Precision: 0.9023
Recall: 0.8992
F1-score: 0.8984
AUC-ROC: 0.9700

Evaluating GBM...

GBM Results:

Accuracy: 0.9412
Precision: 0.9412
Recall: 0.9412
F1-score: 0.9411
AUC-ROC: 0.9786

Evaluating Random Forest...

Random Forest Results:

Accuracy: 0.9412
Precision: 0.9412
Recall: 0.9412
F1-score: 0.9411
AUC-ROC: 0.9691

Summary Table:

	Accuracy	Precision	Recall	F1-score	AUC-ROC
SVM	0.899160	0.902322	0.899160	0.898373	0.970036
GBM	0.941176	0.941202	0.941176	0.941122	0.978597
Random Forest	0.941176	0.941202	0.941176	0.941122	0.969109

Model comparison plot saved as 'model_comparison.png'

Best Parameters:

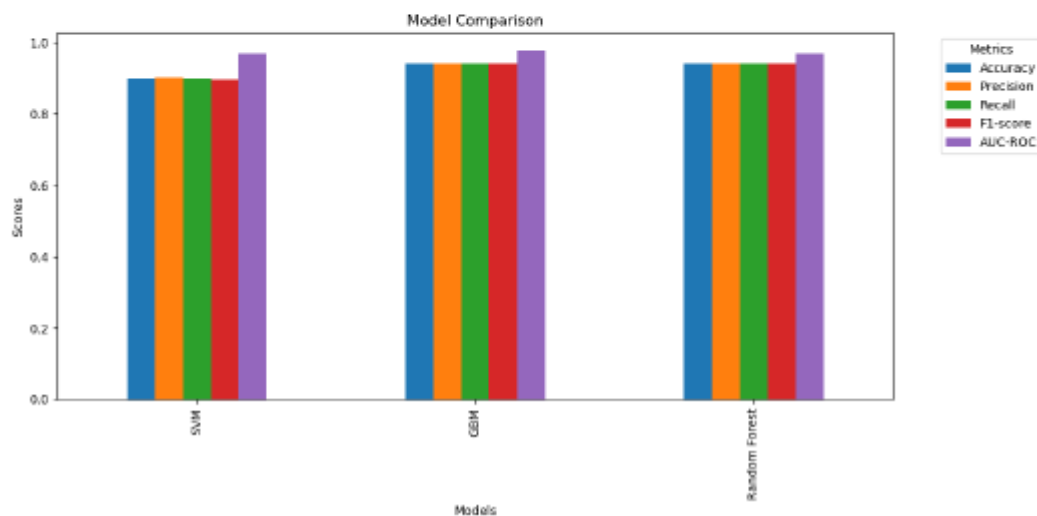
SVM: {'C': 7.390071680409873, 'break_ties': False, 'cache_size': 200, 'class_weight': None, 'coef0': 0.0, 'decision_function_shape': 'ovr', 'degree': 3, 'gamma': 0.7812703466859457, 'kernel': 'rbf', 'max_iter': -1, 'probability': True, 'random_state': 42, 'shrinking': True, 'tol': 0.001, 'verbose': False}

GBM: {'ccp_alpha': 0.0, 'criterion': 'friedman_mse', 'init': None,

```

{'learning_rate': 0.2921569793468812, 'loss': 'log_loss', 'max_depth':
9, 'max_features': None, 'max_leaf_nodes': None,
'min_impurity_decrease': 0.0, 'min_samples_leaf': 2,
'min_samples_split': 9, 'min_weight_fraction_leaf': 0.0,
'n_estimators': 178, 'n_iter_no_change': None, 'random_state': 42,
'subsample': 1.0, 'tol': 0.0001, 'validation_fraction': 0.1,
'verbose': 0, 'warm_start': False}
Random Forest: {'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight':
None, 'criterion': 'gini', 'max_depth': 10, 'max_features': 'auto',
'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease':
0.0, 'min_samples_leaf': 1, 'min_samples_split': 3,
'min_weight_fraction_leaf': 0.0, 'n_estimators': 179, 'n_jobs': None,
'oob_score': False, 'random_state': 42, 'verbose': 0, 'warm_start':
False}

```



Data Preprocessing:

Data is loaded and cleaned. Several preprocessing include data scaling where this was done using StandardScaler to standardize the data.

Model Training:

Three machine learning models are built based on SVM, GBM and Random Forest for comparison. GridSearchCV is used for hyperparameters optimization to decide on the best model parameters.

Model Evaluation:

Once models are trained, they are tested by the use of the test dataset. Evaluations comprising of set measures including accuracy, precision, recall, the F1-score, and the AUC-ROC are determined for each model.

Comparison of Results:

For a compact display, a bar plot is produced to compare the models according to the results of

various evaluation metrics.

Besides, the best parameter for each model is also shown as follows.

Output:

The results of all models are summarized in a table so that a comparison can easily be made.

The model performance comparison consisting of five metrics is shown in the bar chart in the file model_comparison.png.